

Document Validation Report for OCEANAI



Correspondence: Yes

Percentages: 80.0%

Conclusion:

The repository contains implementations of the audio (LSTM, VGG-16, openSMILE), video (MediaPipe FaceMesh, YOLO, ResNet-50), text (Whisper ASR, BERT, LIWC, Bi-LSTM, self-attention), and multimodal fusion (gated fusion, cross-modal transformers) algorithms exactly as described in the documentation. Input formats and key parameters (e.g., log-mel spectrograms, feature dimensions, model architectures) match the technical spec. Pretrained weight downloads and example notebooks demonstrate prediction and accuracy computations. However, dataset loaders, training scripts, and explicit evaluation pipelines with ground-truth performance metrics are not included, limiting full experiment reproducibility.